

F.A.I.L: Financial Anomaly Interpretability using LLMs

MA5741: Object Oriented Programming Course Project

Aritra Dasgupta Ashwini Devharish N Rumaiz Ibrahim K

Department of Mathematics

November 8, 2025

Project Overview & Architecture

The Problem

- Stock anomalies need quick analysis
- Manual review: slow & difficult
- Need actionable insights


Our Solution

- Anomaly detection
- Real-time news
- Historical matching
- AI explanations

Example

Tesla ↓5% → Detect → News → History → Explain

4-Layer Architecture



Key Benefit

Modular design: change one layer independently

Core OOP Concepts & SOLID Principles

Four OOP Pillars

1. Encapsulation

BaseProcessor: hide internals, expose API

2. Inheritance

Processors share logging & progress tracking

3. Polymorphism

IProgressObserver: same interface, different behaviors (Console vs Streamlit)

4. Abstraction

AnomalyDetectionStrategy: defines contract, implementations provide logic

SOLID Principles

✓ Single Responsibility

- DataLoader → load, NewsRetriever → retrieve

✓ Open/Closed

- Add strategies without modifying code

✓ Liskov Substitution

- Swap IRepository implementations

✓ Interface Segregation

- Small, focused interfaces

✓ Dependency Inversion

- Depend on DataLoader, not concrete class

⚠ Real Impact

Makes code testable, maintainable & extensible

Design Patterns & System Workflow

6 Design Patterns

① Strategy

- Runtime algorithm selection
- Data: Yahoo/Alpha — AI: Groq/OpenAI

② Factory

- Centralized object creation

③ Observer

- Progress tracking (Console/Streamlit)

④ Template Method

- Fixed structure, flexible steps

⑤ Builder and Repository

- Fluent config:
`builder.with_ticker().build()`

System Workflow

1. Data Loading (2s)

2. Anomaly Detection

3. News Retrieval

4. Embedding Gen (5s)

5. Similarity Analysis

6. AI Explanation (6s)

i Detection

Z-score $|z| > 2.5 = \text{Outlier} — 17\text{s}$ for 90-day analysis

Real-time progress tracking throughout pipeline

Live System & Deployment



Deployment Status

- **Live on Streamlit Cloud**
- Public access available
- No installation required
- Works on any device



Resources

- Complete source code on GitHub
- Comprehensive documentation
- Setup instructions
- API documentation
- Example use cases



What Makes It Production-Ready?

- ① **Error Handling:** Comprehensive at every layer
- ② **Logging:** Detailed for debugging
- ③ **Fallbacks:** Service failures handled gracefully
- ④ **Configuration:** Easy parameter adjustment
- ⑤ **Modularity:** Independent, testable components
- ⑥ **Documentation:** Clear usage guides



Target Users

- Retail investors
- Financial analysts
- Portfolio managers
- Academic researchers
- Students learning quantitative finance

Technical Achievements & Key Learnings

Metrics

- 3,500+ lines
- 35+ classes
- 6 patterns
- 5 SOLID
- 6 modules
- 17s analysis

OOP Scale

500 lines → messy
3,500 lines → clean

Patterns Work

Strategy: swap providers
Observer: reactive UI



Extensibility

- 1 class = new feature
- Easy provider swap
- Error handling
- Simple testing
- Clear structure

Performance

- Load: 2s
- Embed: 5s
- AI: 6s
- Total:** 17s

Trade-off

More code abstraction
Better long-term maintenance

Contact

[GitHub](#)

[Demo](#)

Takeaway: OOP = Production-ready

Thank You!